

SafeRoute-GNN: A Graph Neural Network Approach for Disaster Response Routing

Authors: Bhavya Jain, Michele Daroowala, Pehar Jhamb, Soumya Pandey

1. Understanding of the Problem

Disaster response routing is critically important for saving lives during emergencies, yet existing navigation systems fail to account for dynamic road blockages caused by floods, earthquakes, landslides, and other disasters.

Key Challenges

- **Dynamic Road Blockages:** Traditional routing systems use static road networks and cannot predict which roads will be blocked during disasters. Manual updates from authorities arrive too late or are incomplete.
- **Disaster-Specific Patterns:** Different disasters (floods vs earthquakes vs landslides) cause fundamentally different blockage patterns. A flood blocks low-lying areas near water, while earthquakes damage bridges and overpasses. Simple distance-based models cannot capture these nuances.
- **Spatial Dependencies:** Road blockage risk depends on complex spatial relationships—proximity to rivers, network topology (degree centrality), terrain elevation, and connectivity patterns that traditional methods cannot model.
- **Limited Labeled Data:** Real disaster data with ground-truth blockages is extremely scarce. Emergency response agencies rarely maintain comprehensive databases of historical road blockages, making supervised learning challenging.
- **Real-Time Requirements:** Emergency routing requires fast inference (seconds, not minutes). Systems must quickly adapt to new disaster scenarios without expensive retraining or manual rule specification.

Core Problem

Existing routing systems **do not explicitly model the spatial dependencies and disaster-specific patterns** that determine road blockages, resulting in routes that guide evacuees into dangerous areas rather than safe shelters.

2. Understanding of Solutions Proposed by Others

2.1 Traditional Shortest Path Algorithms (Dijkstra, A*)

Standard navigation systems (Google Maps, Waze) use shortest path algorithms on static road networks without considering disaster-induced blockages.

Strengths

- Fast and efficient for static networks
- Well-established algorithms with theoretical guarantees
- Optimal path computation on static graphs
- Works well for normal traffic conditions

Limitations

- No awareness of disaster-induced blockages
- Cannot predict which roads will become unsafe
- Requires manual road closure updates from authorities
- Reactive rather than predictive
- Treats all roads uniformly without risk modeling
- No spatial reasoning about disaster propagation

Gap

- No mechanism to predict blockages based on disaster type and spatial features
- Cannot learn from historical patterns or spatial relationships in road networks

2.2 Rule-Based Risk Assessment Systems

Some disaster management systems use hand-crafted rules (e.g., "avoid roads within 500m of rivers during floods") to estimate blockage risk.

Strengths

- Interpretable and transparent decision-making
- Can incorporate domain expert knowledge
- No training data required
- Fast to implement
- Works across different cities with same rules

Limitations

- Cannot capture complex spatial dependencies (e.g., how network topology affects blockage propagation)
- Rules are fixed and do not adapt to new scenarios
- Difficult to scale across multiple disaster types
- Miss subtle patterns (e.g., why certain intersections flood while others nearby don't)
- Threshold selection (500m vs 800m) is arbitrary
- Cannot model interactions between multiple risk factors

Gap

- Manual rules cannot learn from data or discover non-obvious spatial patterns
- No integration of graph structure (connectivity, centrality) with disaster-specific features
- Cannot adapt rules based on observed outcomes

3. Dataset Information and Preparation

For this study, we used real-world road network data from OpenStreetMap (OSM) via the OSMnx library. OSM provides comprehensive, up-to-date street network data for cities worldwide, with rich metadata about road types, infrastructure, and geographic features.

Data Sources

We integrated multiple geographic data sources for comprehensive disaster modeling:

- **Road Network (OSMnx):** Downloaded drive-able road networks for target cities. Includes major roads (highways, arterials) and local streets. Metadata: road length, type (motorway/residential), lanes, surface.
- **Water Bodies:** Rivers, lakes, streams, waterways from OSM natural features. Critical for flood risk assessment. Used to compute proximity-based risk features.
- **Shelter Locations:** Emergency infrastructure from OSM: hospitals, schools, community centers, government buildings, civic centers, social facilities. Target destinations for evacuation routing.
- **Disaster Epicenter Coordinates:** User-specified location or geocoded from city name. Defines the disaster impact zone center. Used to compute distance-based risk features.

Preprocessing Steps

We applied comprehensive graph construction and feature engineering:

1. Downloaded road network for target city/region (Example: Rudraprayag, Uttarakhand, India)
2. Converted OSM MultiDiGraph to undirected graph (bidirectional roads)
3. Extracted node features (6 dimensions per node)
4. Extracted edge features (7 dimensions per edge)
5. Generated ground-truth blockage labels based on disaster simulation
6. Created bidirectional edges for undirected road representation
7. Applied StandardScaler normalization to all features

Dataset Characteristics

Graph Statistics (Example: Rudraprayag, Uttarakhand, India):

- Nodes: 300 road intersections
- Edges: 666 road segments (1332 bidirectional)
- Water features: 352 rivers/streams
- Shelters: 10 emergency locations (hospitals, schools)

Node Features (6D):

8. 1. Latitude - Spatial position (Y coordinate)
9. 2. Longitude - Spatial position (X coordinate)
10. 3. Degree centrality - Number of connected roads (network connectivity)

11. 4. Distance to nearest water body - Flood risk indicator (meters)
12. 5. Distance to disaster epicenter - Proximity to impact zone (meters)
13. 6. Elevation - Terrain feature (simulated: 500-2000m for mountainous regions)

Edge Features (7D):

14. 1. Road length - Segment length in meters
15. 2. Road type encoding - 0 (residential) to 5 (motorway), reflects road importance
16. 3. Number of lanes - Road capacity indicator (1-4+ lanes)
17. 4. Distance to water - Edge-level flood risk (meters)
18. 5. Disaster intensity score - User-specified severity (0-10 scale)
19. 6. Disaster type encoding - Hash-based encoding for ANY disaster type
20. 7. Time since disaster started - Temporal feature (hours elapsed)

Rationale for Choosing This Dataset

- **Authentic Real-World Networks:** Represents actual road topology with realistic connectivity patterns
- **Comprehensive Coverage:** OSM includes roads, water bodies, infrastructure crucial for disaster modeling
- **Global Scalability:** OSMnx works for any city/region worldwide
- **Rich Spatial Features:** Real geographic data enables learning of complex disaster patterns
- **Real-Time Updates:** OSM is continuously updated by mappers
- **Open and Reproducible:** Free, open-source data enables research replication

Label Generation Strategy

Since real disaster blockage data is scarce, we simulate ground-truth labels using disaster-specific heuristics:

- **Flood:** Roads within 500m of water bodies → High blockage risk. Severity-dependent thresholds: Low (250m), Medium (500m), High (1000m).
- **Earthquake:** Random structural damage based on severity probability: Low (10%), Medium (20%), High (35%). Simulates unpredictable infrastructure collapse.
- **Landslide:** Terrain-based blockage simulation (30-50% probability). Higher in mountainous regions.
- **Cyclone/Hurricane:** Wide-area impact (20-55% blockage rate). Severity-dependent.
- **Fire:** Proximity to epicenter with clustering patterns (12-45% blockage).
- **Custom Disasters:** Universal fallback: Proximity to epicenter + intensity-based. Handles ANY user-specified disaster.

Note: In production deployment, these simulated labels would be replaced with historical disaster data from emergency response agencies (e.g., FEMA, Red Cross incident reports).

4. Proposed Method: SafeRoute-GNN

To address the identified gaps, we propose SafeRoute-GNN, which uses Graph Attention Networks (GAT) to learn spatial dependencies in road networks and predict disaster-induced road blockages.

4.1 Graph Representation of Road Networks

We model the road network as a graph $G = (V, E)$ where:

- V = set of intersections (nodes)
- E = set of road segments (edges)
- Node features: $X \in \mathbb{R}^{(N \times 6)}$ - spatial and topological properties
- Edge features: $E_{\text{attr}} \in \mathbb{R}^{(E \times 7)}$ - road characteristics and disaster context

This representation enables the GNN to reason about:

- Topology: How roads connect (centrality, connectivity)
- Spatial proximity: Distances to hazards (water, epicenter)
- Road characteristics: Type, capacity, vulnerability
- Disaster context: Intensity, type, time evolution

4.2 Graph Attention Network Architecture

We use a 3-layer GAT encoder to learn expressive node embeddings:

Layer 1: $\text{GAT}(6D \rightarrow 256D, \text{heads}=4, \text{dropout}=0.3)$

Layer 2: $\text{GAT}(256D \rightarrow 256D, \text{heads}=4, \text{dropout}=0.3)$

Layer 3: $\text{GAT}(256D \rightarrow 128D, \text{heads}=2, \text{dropout}=0.3)$

Why Graph Attention Networks (GAT)?

- Attention Mechanism: Learns which neighboring nodes are important for predicting blockages
- Multi-Head Attention: Multiple attention heads ($4 \rightarrow 4 \rightarrow 2$) capture different types of spatial relationships
- Message Passing: GAT aggregates features from neighbors weighted by learned attention scores

4.3 Edge Blockage Prediction

For each edge (u, v) , we predict blockage probability using an Edge Predictor MLP:

- Step 1: Concatenate source and target node embeddings $\rightarrow [h_u || h_v] \rightarrow 256D$ vector
- Step 2: Feed through MLP: $\text{Linear}(256 \rightarrow 128) + \text{ReLU} + \text{Dropout}(0.3)$, $\text{Linear}(128 \rightarrow 64) + \text{ReLU} + \text{Dropout}(0.2)$, $\text{Linear}(64 \rightarrow 1) + \text{Sigmoid}$
- Step 3: Output blockage probability $\in [0, 1]$

4.4 Route Planning with GNN Predictions

After training, we use the GNN to filter blocked roads and compute safe routes:

21. 1. Input: User location, disaster scenario
22. 2. Predict blockage probability for all edges using trained GNN
23. 3. Threshold filtering: Remove edges where $P(\text{blocked}) > \tau$ (default $\tau=0.5$)
24. 4. Create modified graph G' without high-risk edges
25. 5. Find nearest reachable shelter using Dijkstra on G'
26. 6. Compute shortest safe path from user location to shelter
27. 7. Calculate distance (km) and estimated travel time (minutes)
28. 8. Visualize route on interactive map with blocked roads highlighted

5. Experimentation and Ablation Studies

To validate our approach, we conducted multiple experiments comparing different architectures and configurations.

5.1 Baseline: Standard GNN without Attention

We implemented a baseline using standard Graph Convolutional Networks (GCN) without attention mechanisms.

Results:

- Accuracy: 82%
- F1 Score: 75.3%
- Loss: 0.45

Key Observations:

- GCN learns spatial patterns but treats all neighbors equally
- Missing attention weights reduces ability to focus on critical neighbors
- Performance inferior to GAT-based approach

5.2 Ablation: Removing Disaster-Specific Features

We trained SafeRoute-GNN without disaster-specific edge features (intensity, type, time).

Results:

- Accuracy: 79%
- F1 Score: 71.8%
- Loss: 0.48

Conclusion: Disaster features are crucial. Performance drop of 9% confirms importance of disaster context.

5.3 Ablation: Varying Number of GAT Layers

We experimented with different GAT depths:

Configuration	Accuracy	F1 Score	Loss	Training Time
2 GAT Layers	83%	77.1%	0.43	1.5 min
3 GAT Layers (Ours)	88%	83.3%	0.41	2.2 min
4 GAT Layers	87%	82.5%	0.42	3.8 min
5 GAT Layers	86%	81.9%	0.43	5.1 min

Key Insights:

- 3 layers optimal: Balances expressiveness and efficiency
- Deeper $\neq$ Better: Over-stacking causes over-smoothing
- Diminishing returns: 4-5 layers add complexity without performance gain

6. Results

We compare four configurations on the disaster routing task:

Model	Accuracy	F1 Score	Loss	Blocked Predicted
Baseline (GCN)	82%	75.3%	0.45	95/666
SafeRoute-GNN (No Disaster)	79%	71.8%	0.48	78/666
SafeRoute-GNN (2 Layers)	83%	77.1%	0.43	89/666
SafeRoute-GNN (Final: 3 Layers)	88%	83.3%	0.41	102/666

Example Evaluation Scenario

- Disaster: Earthquake (Medium Severity)
- Location: Rudraprayag, Uttarakhand, India
- Intensity: 8.8/10
- Total road segments: 666
- Predicted blocked: 102 (15.3%)
- Average blockage risk: 24.1%
- Max blockage risk: 87.3%

Routing Results:

- User location: Kedarnath Temple (27.67°N, 85.43°E)
- Nearest safe shelter: Jankalyan Hospital

- Distance: 42.75 km
- Estimated time: 85 minutes (30 km/h avg speed)
- Route segments: 15 road segments
- Safety: Route avoids all high-risk blocked roads

Key Insights

29. 1. GAT Attention Learns Spatial Patterns: Attention weights highlight critical neighboring nodes
30. 2. Disaster-Specific Features Are Critical: 9% performance drop without disaster context
31. 3. 3-Layer Depth Optimal: Balances local and global spatial reasoning
32. 4. High Recall on Blocked Roads: 82% recall means model catches most dangerous roads
33. 5. Efficient Inference: Prediction time <1 second for 666-edge graph

7. Lessons Learnt

1. Graph Structure Matters for Spatial Reasoning

Road networks have rich topology that traditional methods ignore. GNNs explicitly model graph structure, enabling spatial reasoning about disaster propagation.

2. Attention Mechanism Provides Interpretability

GAT attention weights reveal which neighbors influence each prediction, explaining model decisions.

3. Domain-Specific Features Are Non-Negotiable

Removing disaster features caused 9% accuracy drop. Spatial features alone cannot distinguish disaster types.

4. Simulated Labels Work for Proof-of-Concept

Heuristic-based label generation enabled model development without expensive real data collection.

5. Deeper Is Not Always Better

Over-stacking GAT layers (4-5) caused over-smoothing. 3 layers balanced local and global reasoning.

6. Real-World Deployment Requires Simplicity

Multi-model ensembles achieved +2.5% accuracy but 3× cost. Single GAT model is more practical.

7. Ablation Studies Reveal What Matters

Systematic ablations identified critical components: attention (+6%), disaster features (+9%), 3-layer depth (+5%).

8. Open-Source Data Enables Global Scalability

Using OSMnx means SafeRoute-GNN works for any city worldwide without custom data collection.

8. Future Work and Limitations

Limitations

- Simulated Labels: Ground truth generated via heuristics, not real disaster data
- Static Disaster Assumptions: Model assumes disaster characteristics don't evolve during evacuation
- No Traffic Congestion: Doesn't account for evacuation traffic
- Limited Validation: Tested on limited scenarios, need multi-city evaluation
- No Uncertainty Quantification: Outputs point probabilities, not confidence intervals

Future Directions

- Real Disaster Data Integration: Partner with emergency agencies for historical blockage data
- Temporal Graph Neural Networks: Model disaster evolution over time
- Multi-Agent Routing: Account for congestion from multiple evacuees
- Satellite Imagery Integration: Use real-time imagery to detect actual blockages
- Mobile Application: Deploy as mobile app for citizens
- Multi-City Benchmarking: Test on diverse geographies
- Explainability Dashboard: Visualize attention weights for transparency

9. References

1. Graph Neural Networks

Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P., & Bengio, Y. (2018). Graph Attention Networks. International Conference on Learning Representations (ICLR).

Kipf, T. N., & Welling, M. (2017). Semi-Supervised Classification with Graph Convolutional Networks. ICLR.

Hamilton, W. L., Ying, R., & Leskovec, J. (2017). Inductive Representation Learning on Large Graphs. NeurIPS.

2. Disaster Response and Routing

Litman, T. (2006). Lessons from Katrina and Rita: What Major Disasters Can Teach Transportation Planners. Journal of Transportation Engineering, 132(1), 11-18.

Cova, T. J., & Johnson, J. P. (2003). A Network Flow Model for Lane-Based Evacuation Routing. Transportation Research Part A, 37(7), 579-604.

3. Geospatial Data and Tools

Boeing, G. (2017). OSMnx: New Methods for Acquiring, Constructing, Analyzing, and Visualizing Complex Street Networks. Computers, Environment and Urban Systems, 65, 126-139.

OpenStreetMap Contributors. (2024). OpenStreetMap. <https://www.openstreetmap.org>

4. PyTorch and Graph Libraries

Fey, M., & Lenssen, J. E. (2019). Fast Graph Representation Learning with PyTorch Geometric. ICLR Workshop.

Paszke, A., Gross, S., Massa, F., et al. (2019). PyTorch: An Imperative Style, High-Performance Deep Learning Library. NeurIPS.

10. Conclusion

SafeRoute-GNN demonstrates that Graph Neural Networks can effectively learn spatial dependencies in road networks to predict disaster-induced blockages. By combining:

- Graph Attention Networks for spatial reasoning
- Disaster-specific features for context awareness
- Real-world OSM data for authentic network topology
- Efficient inference for real-time emergency response

we achieve 88% accuracy in predicting road blockages, enabling safe evacuation routing that avoids dangerous areas.

The system is globally scalable (works for any city), interpretable (attention weights explain predictions), and practical for deployment in emergency management systems.

Impact: SafeRoute-GNN can save lives by guiding evacuees away from blocked roads during floods, earthquakes, and other disasters.